

Machine Learning

Linear Regression

Homework 2

Mostafa S. Ibrahim

Teaching, Training and Coaching for more than a decade!

Artificial Intelligence & Computer Vision Researcher

PhD from Simon Fraser University - Canada

Bachelor / MSc from Cairo University - Egypt

Ex-(Software Engineer / ICPC World Finalist)



© 2023 All rights reserved.

Please do not reproduce or redistribute this work without permission from the author

Tips for this assignment

- Overall, this is an easy assignment
- However, it consists of many requests, and that can be overwhelming!
- My advice:
 - First, see all the requirements
 - Start with the first task, because this is the CORE of the assignment. Everything else depends on it

Use ArgumentParser

- It is important to have a good coding that allows easy experimentation
- Use Python's ArgumentParser

```
parser = argparse.ArgumentParser(description='Linear Regression Homework')

parser.add_argument('--dataset', type=str, default='dataset_200x4_regression.csv')

parser.add_argument('--preprocessing', type=int, default=1,          # p4
                    help='0 for no processing, '
                         '1 for min/max scaling and '
                         '2 for standrizing')

parser.add_argument('--choice', type=int, default=2,
                    help='0 for linear verification, '          # p0
                         '1 for training with all features, '    # p1 / p3 / p7
                         '2 for training with the best feature, ' # p5
                         '3 for normal equations, '              # p6
                         '4 for sikit')

# Use below to explore for p2
parser.add_argument('--step_size', type=float, default=0.01, help='Learning rate (default: 0.01)')
parser.add_argument('--precision', type=float, default=0.0001, help='Requested precision (default: 0.0001)')
parser.add_argument('--max_iter', type=int, default=10000, help='number of epochs to train (default: 1000)')

args = parser.parse_args()
```

P1: Implement LR using Gradient Descent

- We already coded the GD algorithm for function minimization
- Now, we just add a **few lines of code** for **linear regression**
- We just need to code the **derivative of MSE** over N examples
 - Try to do that in a Pythonic way (vectorization - *2-4 lines of code*)
- **def** gradient_descent_linear_regression(X, t, step_size = 0.01, precision = 0.0001, max_iter = 1000):
 - X is an N * D dimensions input vector (first column with ones for the intercept)

$$cost(W) = \frac{1}{2N} \sum_{n=1}^N (y(X^n, W) - t^n)^2$$


$$\frac{\partial cost(W)}{\partial W_j} = \frac{1}{N} \sum_{n=1}^N (y(X^n, W) - t^n) * X_j^n$$

Early Verifications

- You should always think how to verify your code
 - I take baby steps to compile, build and verify
- The easiest way is to start with a simple data for a **line (no noise)**
 - E.g. a 45 degree line
 - You should be able to perfectly fit it
- Consider the following data:
 - `x = np.array([0, 0.2, 0.4, 0.8, 1.0])`
 - `t = 5 + x`
 - Clearly the solution for such data is: slope = 1 and intercept = 5
- We already implemented its cost function and derivative

Visual Verification

- Draw your **found** line. It should perfectly fit the data
- I used the following configuration:
 - `step_size=0.1`, `precision = 0.00001`, `max_iter=1000`
- Start your weights from a random position
 - `cur_weights = np.random.rand(features)`
- Note: try other learning rates (step size)
 - E.g. 0.01 0.001



Dataset

- Now, you verified the algorithm. Time to work on the requested dataset
- Attached a simple dataset: **dataset_200x4_regression.csv**
- It consists of 200 x 4 matrix
 - The first 3 columns represents **3 features** (x1, x2, x3)
 - The last column represents **our target** that we want to regress it
- The data is numeric and very clean
 - The real datasets has non-numeric values (e.g. city name)
 - The real datasets are not clean (e.g. missing values)
- Start with using MinMax Scaler for the data
- We will use the 200 examples for training

More Verifications

- Use the following information for the verification step
 - Process all the data using **min-max scaler**
 - Start from **initial state: [1, 1, 1, 1]**
 - 3 input features \Rightarrow learn 4 parameters (includes intercept)
 - Use **step_size = 0.1 and max_iter = 3**


```
weights: [1. 1. 1. 1.]
cost: 1.5127403213757413 - gradient: [1.69732353 0.87760002 0.88852874 0.50521329]
weights: [0.83026767 0.91223997 0.9111471 0.9494787 ]
cost: 1.0800099415415019 - gradient: [1.42904069 0.73688578 0.75333759 0.42991259]
weights: [0.6873636 0.8385514 0.83581334 0.9064874 ]
cost: 0.7724470467820816 - gradient: [1.20289468 0.61829309 0.63933402 0.36641601]
```


P2: Optimizing the **hyper**parameters

- Let's fix the maximum number of iterations to 10,000
- Consider the following 2 hyperparameter sets:
 - Step sizes: {0.1, 0.01, 0.001, 0.0001, 0.00001, 0.0000001}
 - Precision: {0.01, 0.001, 0.0001, 0.00001}
 - For each combination, run 3 times of your program to try different **initial** startings
- For all the possible setups, investigate:
 - What is the minimum error you achieved among them?
 - How many iterations before the program stops?
 - What is your best combination?

P3: Visualization

- Visualize the iterations versus the cost function of each state your program passed with
- How can such a graph help us investigate the performance?




P4: Exploration

- What will happen if we did not preprocess the data (e.g. scaling)?
- What about standardizing the data (e.g. scikit StandardScaler)?

P5: Data **Linearity** Investigation

- So far, we just used the whole data as they are
- When we started linear regression, we assumed data seems coming from a line
- Is that true?
- Draw 3 plots, each one is a **feature versus the target**
 - Which feature seems really came from a line?
- Retrain your **model only using this feature**
($mx+c$ case)



P6: Normal Equations

- **The Normal equation** is a closed-form solution (OLS) for the linear regression problem
- Implement function: `def normal_equations_solution(X, y):` $\Theta = (X^T X)^{-1} X^T y$
- Given the data, the function computes the parameter
- Compute the results **with and without scaling**
- Compute the actual predictions
- Any thoughts?
- There are 2 ways to implement this expression. Can you find the most efficient way?

My results

- Let me share some results for the sake of verification
- Let's fix: $\text{step_size} = 0.01$, $\text{precision} = 0.0001$, $\text{max_iter} = 1000$
- Scaling + Gradient Descent + 3 features
 - *Number of iterations ended at **2021** - with cost 0.0033275517970179813 -*
 - *Optimal weights [0.15874904 0.54368119 0.10489429 0.22107806]*
 - *The actual cost function on the original domain is **116.5076769533999***
- Scaling + Normal Equations + 3 features
 - *Optimal weights: [0.12060381 0.6338432 0.20894728 0.00150253]*
 - *The actual cost function on the original domain is **118.11925272653018***


P7: Gradient Check

- *This question comes in **interviews***
- Luckily, in this homework I give you output samples so that you can **verify your derivative code**
- In practice, you might compute the derivative wrongly. Or compute it correctly and code it wrongly
- Luckily, in programming, there's a way to check if we coded the derivative correctly or not - using a simple formula for the gradient!
- Tip: this feature is implemented in the [SciPy](#) library

Gradient Check

- From calculus, we know we can compute the **gradient numerically** from the cost function (let's call it J)
- What does this formula say?
- Consider a segment of length $2C$ centred around the current weights

$$\frac{\partial}{\partial \theta_i} J(\theta) \approx \frac{J(\theta_{i+c}) - J(\theta_{i-c})}{2c}$$



Gradient Check

- It says if you want to compute the derivative of your function with respect to the i -th variable, you should consider the following:
 - Use a small epsilon variable C , e.g. $= 1e-4$
 - Compute the cost function at 2 locations: $\text{weight}[i] + c$ and $\text{weight}[i] - c$ (other weights fixed)
 - Compute the above formula. This is the derivative estimate for the i -th weight
 - Verify this gradient is the same as your gradient from the partial derivatives
 - Do this for every variable
- Enhance your gradient code with a **gradient check**
- Question: why not simply employ this gradient alternative in practice?

Code Organization

- In real projects, we don't just write a lot of code in the same file
- We need to partition the code to files, packages, classes and functions
- Think how you will structure your code
- Here is mine

Name	Size	Type
data	1 item	Folder
assignment_driver.py	3.1 kB	Text
data_helper.py	856 bytes	Text
gradient_descent_linear_reg.py	1.7 kB	Text
normal_eq.py	158 bytes	Text
regression_utilites.py	990 bytes	Text

P8: Simple Linear Regression Derivation

- We mentioned the 2D line fitting $y=mx+c$ has a simple formula. In this task, you will start from the cost function and derive the formula
- Let our line formula be: $\hat{Y}_i = a + Bx_i$ (here, a is like c , and B like m)
- Our cost function is
$$S = \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$
- To find the minimum:
compute derivative = 0
Then extract the target variable

Simple Linear Regression Derivation : Steps

- First compute the partial derivative dS/da
- You should end with the formula:
 - The bar means the mean(Y)
- Then compute dS/dB

$$a = \bar{Y} - B\bar{x}$$

- Substitute a with the formula we derived
- Rearrange to end with the formula

$$B = \frac{\sum_{i=1}^n (x_i Y_i - \bar{Y} x_i)}{\sum_{i=1}^n (x_i^2 - \bar{x} x_i)}$$

P9: Time Complexity

- Assume that:
 - k is number of iterations
 - n is number of training examples
 - d is number of features
- Analyze the **time complexity** for
 - Batch Gradient Descent Algorithm
 - Normal Equation
 - Compare them
- Recall: Time complexity for matrix multiplication
 - $[A \times B] * [B \times C]$ is $O(ABC)$. If $A=B=C$, then is $O(A^3)$

“Acquire knowledge and impart it to the people.”

“Seek knowledge from the Cradle to the Grave.”

