

Machine Learning

Boosting Model Inference Speed

Mostafa S. Ibrahim

Teaching, Training and Coaching for more than a decade!

Artificial Intelligence & Computer Vision Researcher

PhD from Simon Fraser University - Canada

Bachelor / MSc from Cairo University - Egypt

Ex-(Software Engineer / ICPC World Finalist)



© 2023 All rights reserved.


Please do not reproduce or redistribute this work without permission from the author

Operations speed

- FLOPS stands for "Floating Point Operations Per Second", a common measure to calculate the number of floating-point calculations
 - GFLOPS (Giga), TFLOPS (trillions), PFLOPS (quadrillions)
- However, still there are many issues with the metric
 - Companies vary on what is an operation (e.g. fusing 2 operations)
 - Apps typically have lower utilization (job's flops / available flops)
 - Doesn't take into considerations: memory bandwidth, I/O operations / Integer operations
- MLPerf Benchmarks: [Nvidia](#) / [general](#)

Boosting Model Inference

- Running machine learning models requires **extensive computations and memory**
 - It is cost-saving to have so **fast** models but still with **high** performance
 - Tip: in startups we may skip and validate the idea first (MVP)
 - In real-time apps (e.g. self-driving), you must be extremely fast!
- Imagine we process a live stream for basketball Olympiad
 - Our ML system should detect **winning points!**




Boosting Model Inference

- There are 2 stages that we do to speed up a model's **inference**!
- 1) Model Compression
 - **First**, think of a simpler network that still has good performance (architecture optimization)
 - Neural Architecture Search (NAS) explores better architectures
 - Convolution **factorization**: $1 \times N$, $N \times 1$, 1×1 , depthwise conv (low-rank factorization)
 - **Second**, use Techniques to reduce the model size of the final network
 - Cons: You may **lose** some **accuracy**
- 2) Inference Optimization
 - **Software and hardware** are very specialized to **accelerate** the network inference
- In a real-time app like self-driving, we do all of them!

Model Compression: Simpler Model

- Imagine your first model reads a 3-seconds video at 60 FPS
 - The model requires **157 Million** parameter
- You then realized at 60 FPS, many frames are duplicate
 - then you sampled only 5 frames per second
- You also realized a simpler backbone (like MobileNetV2) works well
 - And with some other tricks, you reduced the network into **13M** parameters!

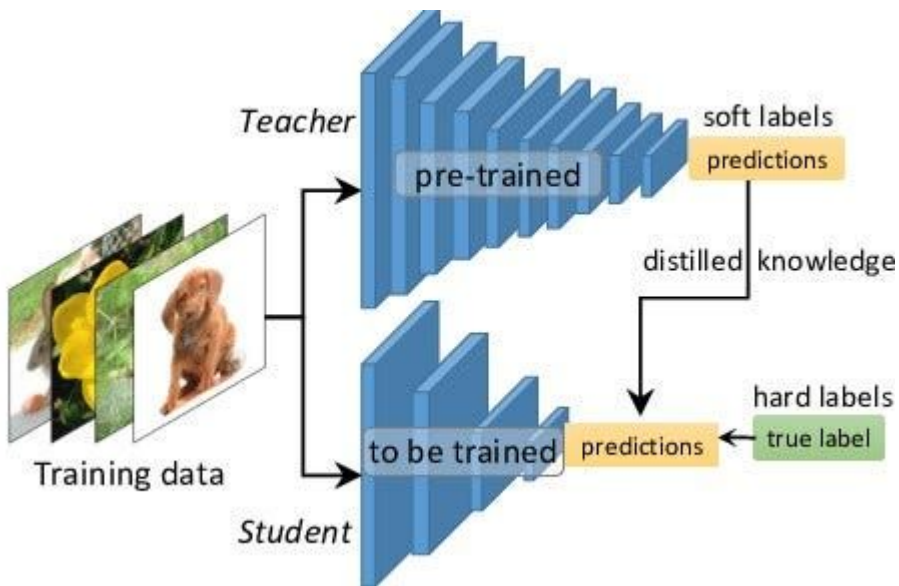


Model Compression: Common Techniques

- Research is so active on approaches to reduce the network size
- Some of these approaches are easily applicable, while others may need more development efforts
- Common techniques
 - **Knowledge Distillation:** a smaller **student** DNN model is trained to [transfer the knowledge](#) of a larger **teacher** DNN (nice but requires effort to train the teacher and transfer)
 - **Pruning:** Pruning removes the neurons or connections in the neural network that contribute the least to the model's performance (Filter/Layer Pruning - Sparse/Dynamic Pruning)
 - Also can be applied to some other algorithms
 - **Quantization:** very general/common in industry as easily applicable in many cases
 - E.g. convert network from 32-bit **floats** to 8-bit **integers**.

Knowledge Distillation

- **Knowledge Distillation:** a smaller **student** DNN model is trained to **transfer the knowledge** of a larger **teacher** DNN (nice but requires effort to train the teacher and transfer)



Model Compression: Common Techniques

0.34	3.75	5.64
1.12	2.7	-0.9
-4.7	0.68	1.43


FP32



Quantization


64	134	217
76	119	21
3	81	99

INT8



pruning
synapses

pruning
neurons



Model Compression: Quantization

- Quantization **reduces the precision** of the numerical values in the model, converting **32-bit** floats to 16-bit or even **8-bit** integers.
 - Significantly reduce the model size but with a slight degradation in performance.
 - **8-bit** is very common in the industry. Binary quantization is limited
- It is supported by Pytorch and Tensorflow frameworks
 - **Quantization-Aware Training (QAT)**
 - Most common. First, train your normal F32 bit float network
 - Retrain it: Fine-tune it 2-5 epochs with quantization
 - Post-Training Quantization (PTQ)
 - The floating model is quantized after the training (no re-training)
 - Relevant: **Layer fusion**: an optimization technique that combines multiple adjacent layers in a neural network into a single layer

Model Compression: Evaluation

- Given that we have several ways to compress a model, it is good to think in how to criteria to compare
- **Dependency:** Each compression technique exhibits different dependencies on the model's architecture and the target hardware
 - Which technique is DNN agnostic?
 - Which technique is Hardware independent?
- **Speed Gains**
 - Which one of them is faster?
- **Tip:** nothing prevents us from applying all of them
 - For example, do distillation first to reduce the network
 - Then do pruning to remove useless layers/connections
 - Then Quantize the weights!

Model Compression: Dependency

- **Quantization**

- Hardware **dependent** (e.g. bits/ram), but **independent** approach for most of DNNs
- Deeper networks might tolerate aggressive quantization better than shallower ones.

- **Pruning**

- Different architectures might require specific pruning strategies (**dependent**)
- **Independent Hardware:** Even with a lot of pruning, CPUs/GPUs are not optimized for **sparse** operations (RAM also degrade performance due to unstructured pruning)
 - Built for dense not sparse operations
- **Dependent Hardware:** developed to handle sparse computations efficiently

- **Distillation**

- Student network design is important to learn efficiently / capture the [dark knowledge](#)
- Hardware **independent**. Select a student network that fits with required constraints

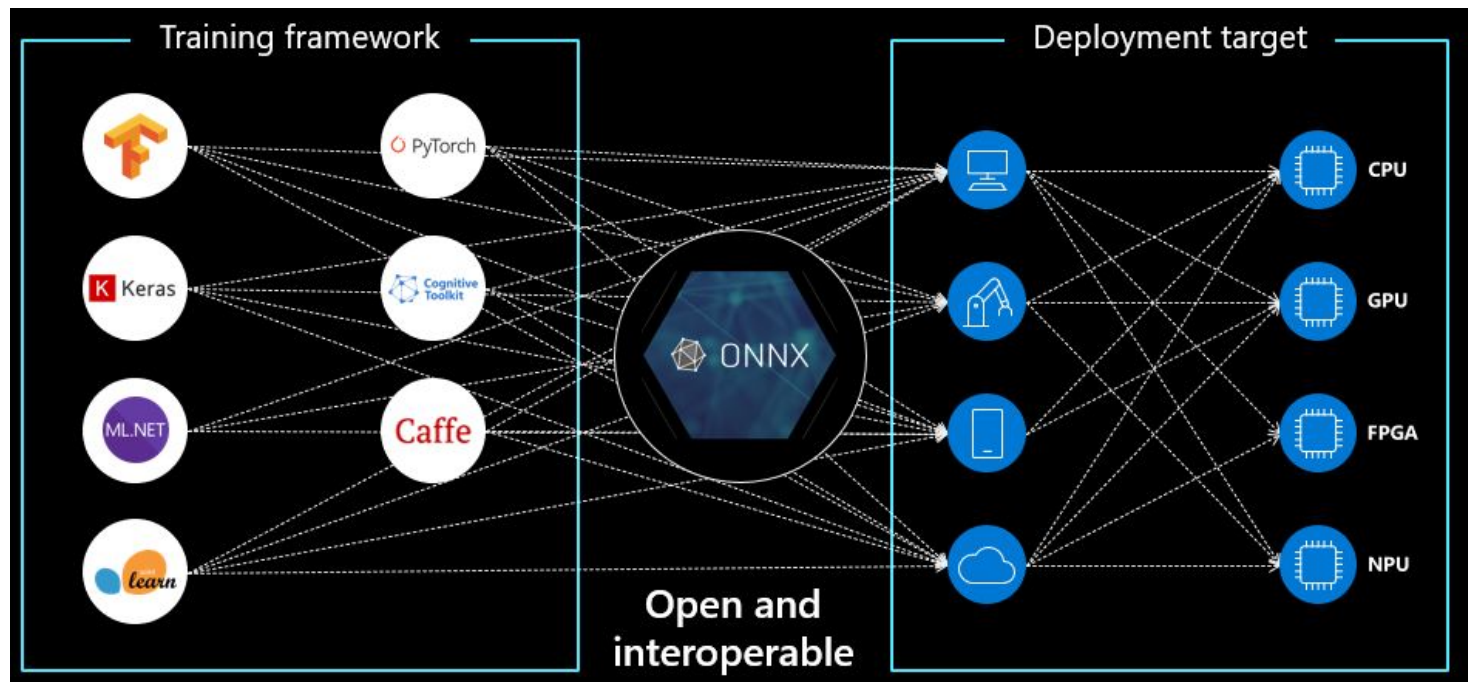
Model Compression: Speed Gains

- **Quantization** can bring 2-4x speed.
 - 8-bits with quantization-aware training is very common choice
 - Lower than will cause performance drop
- **Pruning**
 - The speed gain depends on the 1) sparsity level and the 2) hardware's ability
 - General-purpose hardware might not see a significant speedup
 - Specialized hardware / software may achieve speedups
- **Distillation**
 - The speed gain depends depends on the size of the student models compare to the teacher

Inference Optimization: ONNX Role

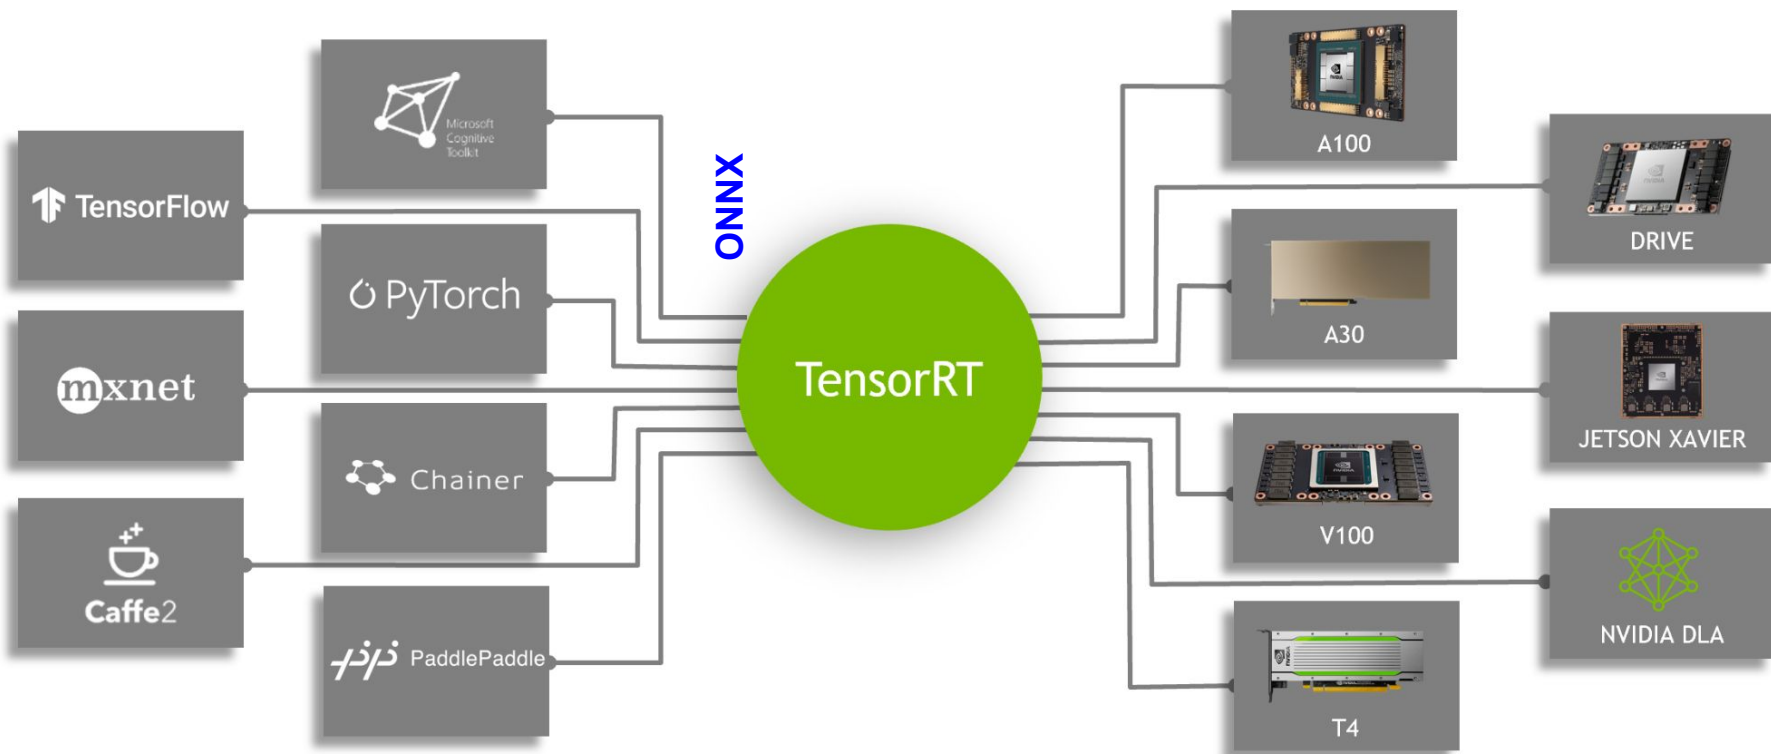
- ONNX library convert models from different ML libraries (e.g. PyTorch, TensorFlow) into a unified format (ONNX format) ⇒ **interoperability**
- This way, software/hardware **inference optimization** techniques focus only on optimizing ONNX files toward different deployment targets (e.g., CPU, GPU)
- Challenges
 - Version **Compatibility**
 - Operator Support: Not all layers and operations in every framework are supported
 - **We typically change the backbone to use the supported operators**
 - Harder to debug the model searching for conversion issues

Inference Optimization: ONNX



NVIDIA: Optimized Software and Hardware

- NVIDIA TensorRT is a **high-performance** deep learning **inference library** for deployment of neural network models. It is designed to work efficiently on **NVIDIA** GPUs and is part of NVIDIA's Deep Learning SDK
 - One major input files are the ONNX files. It optimizes them further
- **[NVIDIA's] DLA** (Deep Learning Accelerator) refers to **specialized hardware** designed to **accelerate** deep learning tasks, particularly **inference**
 - Primarily used for inference in edge devices like cameras, smartphones, and IoT devices.
 - Optimized for lower latency and **power consumption**
 - Note: GPUs can be used for both **train and inference**



PyTorch Or Tensorflow Or Caffe Library: **Train and Export**

Network
Training

Trained
Model
32-bit

Quantize
QAT/PTQ

Quantized
Model
8-bit

export

ONNX
Model
8-bit

Dataset

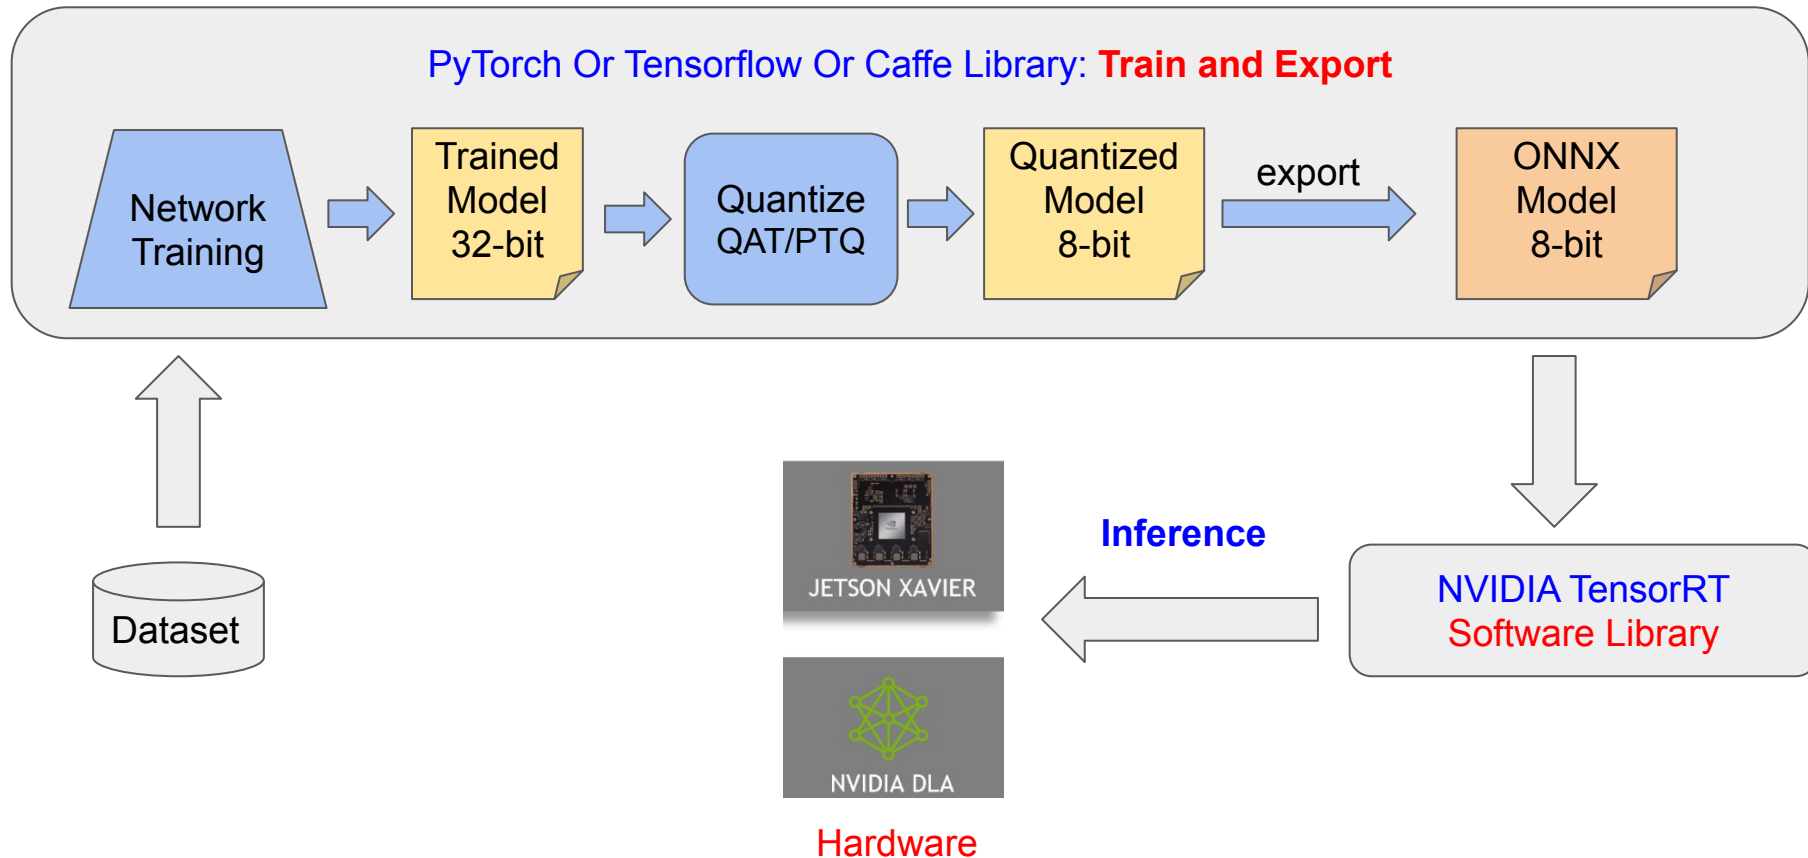
Inference

NVIDIA TensorRT
Software Library

JETSON XAVIER

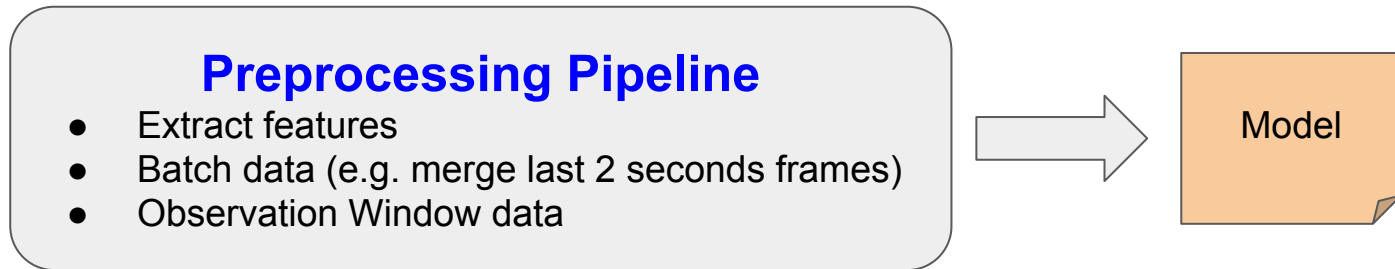
NVIDIA DLA

Hardware



From Python (Development) to C++ (Deployment)

- Nowadays, we typically do DNN **training** using python with frameworks such as PyTorch and Tensorflow
- However, in several applications we run the **inference** on a C++ pipeline to speed up (e.g. in embedded systems, real-time like self-driving)
- One of the common **mistakes** is failing to **sync code** changes between **training python** pipeline and **inference C++** pipeline
 - Wrong Code Logic (*could be 2 teams responsibilities*)
 - Changes in configurations (e.g. temporal window length)



Offline Inference

- Sometimes our model just runs offline, which allows different things
 - For example, google process offline your images and extract information
- We can boost the performance in 2 ways:
 - Model Ensemble (average results of multiple models)
 - Test Time Augmentation ([TTA](#)): Average results of multiple argumentations
 - In real inference, we just go with the input as it is

“Acquire knowledge and impart it to the people.”

“Seek knowledge from the Cradle to the Grave.”