

# Machine Learning

## Probabilistic Modeling

**Mostafa S. Ibrahim**


*Teaching, Training and Coaching for more than a decade!*

*Artificial Intelligence & Computer Vision Researcher*

*PhD from Simon Fraser University - Canada*

*Bachelor / MSc from Cairo University - Egypt*

*Ex-(Software Engineer / ICPC World Finalist)*

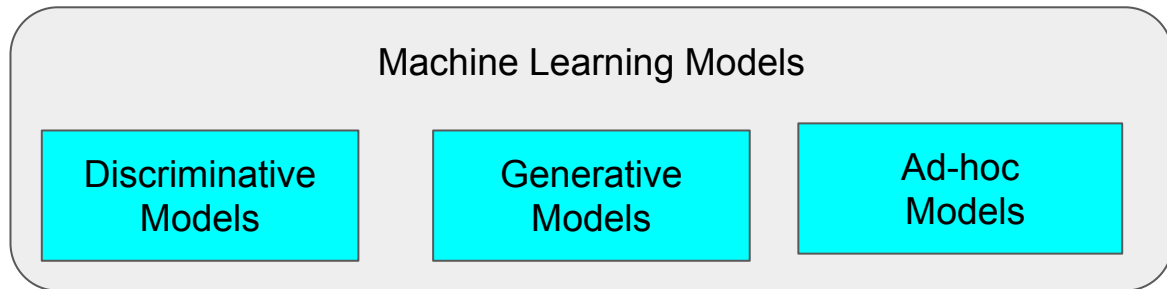


© 2023 All rights reserved.


Please do not reproduce or redistribute this work without permission from the author

# Keep in mind

- This is a **high level** lecture to show you another perspective in ML
  - No intention to dig in any of these details
  - It will be part of your later journey
- Most of ML models are either **Discriminative** (this course) or **Generative** Models (probabilistic nature).
  - A few models are neither (e.g. **K-means**, learning policy in reinforcement learning and some anomaly detection algorithms, etc)



# Recall: Probability



Joint Probability

$$P(X = x_i, Y = y_j) = \frac{n_{i,j}}{N}$$


Marginal Probability

$$P(X = x_i) = \frac{c_i}{N}$$

Conditional Probability

$$P(Y = y_j | X = x_i) = \frac{n_{i,j}}{c_i}$$

# Recall: Probability



Marginal Probability

$$P(X = x_i) = \frac{c_i}{N}$$


Conditional Probability

$$P(Y = y_j | X = x_i) = \frac{n_{i,j}}{c_i}$$

Product Rule

$$\begin{aligned} P(X = x_i, Y = y_j) &= \frac{n_{i,j}}{N} = \frac{n_{i,j}}{c_i} \cdot \frac{c_i}{N} \\ &= P(Y = y_j | X = x_i) P(X = x_i) \end{aligned}$$

# Recall: Probability



Marginal Probability

$$P(X = x_i) = \frac{c_i}{N}$$

Conditional Probability

$$P(Y = y_j | X = x_i) = \frac{n_{i,j}}{c_i}$$


Product Rule

$$\begin{aligned} P(X = x_i, Y = y_j) &= \frac{n_{i,j}}{N} = \frac{n_{i,j}}{c_i} \cdot \frac{c_i}{N} \\ &= P(Y = y_j | X = x_i) P(X = x_i) \end{aligned}$$

Sum Rule	$p(x) = \sum_y p(x, y)$
Product Rule	$p(x, y) = p(y   x) p(x)$

# Recall: Probability Density Function

- A probability density function of continuous random variable, is a function whose value at any given sample in the sample space can be interpreted as providing a relative **likelihood** that the value of the random variable would be equal to that sample
- Below the Gaussian PDF



$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$


# Question!

- Given a Gaussian distribution with mean 1 and sigma 0.5 and a sample  $x = 0.25$ . How likely that sample **belongs to (sampled from)** this distribution?
- Simply evaluate the pdf  $N(x=0.25 \mid \text{mean}=1, \text{std}=0.5^2) \approx 0.0841$ 
  - This is not a probability, but a density.
  - The higher the density, the more likely it is that  $x=0.25$  was sampled from this Gaussian distribution.

$$f(0.25|1, 0.5^2) = \frac{1}{\sqrt{2\pi(0.5)^2}} \exp\left(-\frac{(0.25 - 1)^2}{2(0.5)^2}\right)$$

# Question!

- Below is a gaussian  $p(x|\theta)$  of students weights
  - where  $\theta$  is gaussian mean and sigma
- Given 3 **independent** samples,  $x_1$ ,  $x_2$ ,  $x_3$
- How likely all of them came from this distribution?
- $p(x_1|\theta) * p(x_2|\theta) * p(x_3|\theta)$








# Likelihood

- Assume we have a dataset of N **independent** examples:  $X_1, X_2, \dots, X_n$ 
  - We know the data comes from some distribution, e.g. gaussian
- Given **parameters** (weights)  $\theta$  of a distribution, what is the probability of observing the dataset?
  - $\Theta$  is fixed. Evaluate the data [reasoning context]
  - Intuitively, just multiply  $p(x)$  of all the values!
- Another perspective  $L(\theta|\text{Data})$  - *Optimization context*
  - Given a dataset, how "likely" are different values of  $\theta$ ?
    - Likelihood of the parameters  $\theta$ , given the observed data
  - Data is fixed. Try different  $\theta$  to pick the best  $\theta$  using MLE (Optimization context)
  - *Not a probability function: integrate over different  $\theta \neq 1$*

$$\mathcal{L}(\theta|X) = P(X|\theta) = \prod_{i=1}^n P(x_i|\theta)$$

# Question!


- We have a dataset of 1000 samples below in green from unknown gaussian
- Which gaussian seems the one that used to sample them? A, D or F
- F. When we multiply all  $p(x)$ , it will be higher than the others!
  - **We expect most of the data to be at the mean + symmetry**
  - Little problem. Some value might be  $p(x) = 0$ ! Hence likelihood = 0!



# From Likelihood to Log-Likelihood


- Logarithm is **monotonic**, hence we can replace the multiplication of probabilities with the sum of **logs of probabilities**
  - We don't care about the exact values. Only relative value to **compare**

$$P(X|\theta) = \prod_{i=1}^n (p(x_i|\mu, \sigma^2))$$



$$\ell(\mu, \sigma^2) = \sum_{i=1}^n \log (p(x_i|\mu, \sigma^2))$$

- Recall:  $\log(a*b) = \log(a) + \log(b)$



# From Likelihood to Log-Likelihood

- We can use the **log-likelihood score** as an indicator for the likelihood data given a distribution
  - Below X is a an array

```
def gaussian(x, mu, sigma):  
    return (1 / (np.sqrt(2 * np.pi * sigma ** 2))) * \  
        np.exp(-0.5 * ((x - mu) / sigma) ** 2)
```

```
def likelihood_gaussian(data, mu, sigma):  
    return np.prod(gaussian(data, mu, sigma))
```


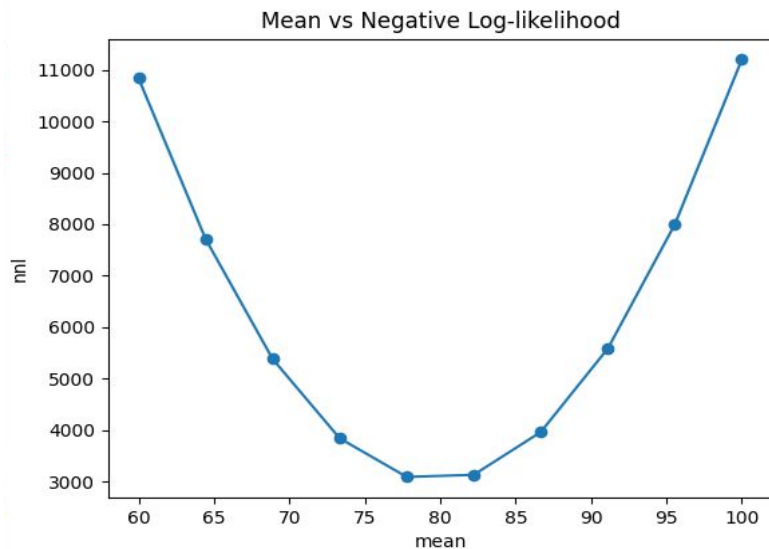
```
def log_likelihood_gaussian(data, mu, sigma):  
    return np.sum(np.log(gaussian(data, mu, sigma)))
```

$$\prod_{i=1}^n (p(x_i | \mu, \sigma^2))$$

$$\sum_{i=1}^n \log (p(x_i | \mu, \sigma^2))$$

# Maximum Likelihood Estimation (MLE)

- The best gaussian representing the data is the one with the maximum likelihood. Let's try means in range [60-100] with fixed sigma=5
  - Maximum likelihood is the same indicator as maximum log-likelihood
  - This is also the same as the **minimum negative log-likelihood**, as  $\text{Max } F = - \text{Min } F$



# Maximum Likelihood Estimation (MLE)

- MLE is the *most common* approach to find the optimal parameters (weights) to **fit a distribution** (e.g. a gaussian) to a given data coming from unknown distribution
  - MLE find the parameter **values** that **maximize** the likelihood function
- How to implement
  - First decide the probability function P
    - E.g. bernoulli or gaussian
  - Second, apply  $F = -\log(P)$
  - Write down the derivative of F
  - Now, you have a function and its derivative, just apply **gradient descent** to find the minimum
    - The minimum represents the **best parameter estimation**
  - Play in the homework

$$\mathcal{L}(\theta|X) = P(X|\theta) = \prod_{i=1}^n P(x_i|\theta)$$

# Why prefer log probabilities?!

- Numerical Stability
  - very small numbers  $\Rightarrow$  underflow / very large numbers  $\Rightarrow$  overflow
- Simplification
  - Multiplication to Addition:  $\log(a*b) = \log(a) + \log(b)$ 
    - Easier to model and compute, e.g. in likelihood
  - Division to Subtraction
- Differentiability: Gradient Descent friendly
- Interpretability
  - Exponential Families: Log probabilities yield simpler forms for these distributions
  - Log-odds e.g. in logistic regression
  - Information Theory: Logarithms are the basis for entropy and information measures

# MLE for Univariate Gaussian

- Given a dataset of a **univariate**  $X$  coming from a **Gaussian** distribution, the MLE is just the sample mean and variance of  $X$ 
  - We can compute that analytically [see homework]
  - Maximum likelihood ignores parameter uncertainty (think of a single example)
- Biased sample variance issue
  - The finite sample mean is biased to the sample. But we know **population mean** is different
  - The bigger problem, the dataset variance is based on the biased dataset mean
  - There are approaches to reduce this bias due to finite sample size (e.g. **Bessel's correction**)
    - Divide by  $n-1$  instead of  $n$  to convert biased variance to **unbiased** sample variance

Unbiased sample variance

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

MLE variance (biased)

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

Convert MLE to unbiased

$$s^2 = \left( \frac{n}{n-1} \right) \sigma^2$$



# Equivalence

- We can prove that **minimizing the cross-entropy** is equivalent to:
  - Minimizing KL divergence [homework]
  - **Maximizing the likelihood** = Maximizing the **log** likelihood
    - Why? the logarithm is monotonic
  - Maximizing the **log** likelihood = Minimizing the **negative log**-likelihood
    - Why? maximizing  $F$  = minimizing  $-F$

# Equivalence

- Assume a dataset of  $N$  examples over  $C$  classes for a multi-classifier
  - The likelihood function of **observing** the given set of **labels** given the **predictions**
    - Labels  $y$  are typically one-hot-encoding
  - Apply the -log
    - Switch 2 multiplications into 2 sums
    - Convert internal to  $y_i \log y_i$
    - Add negative

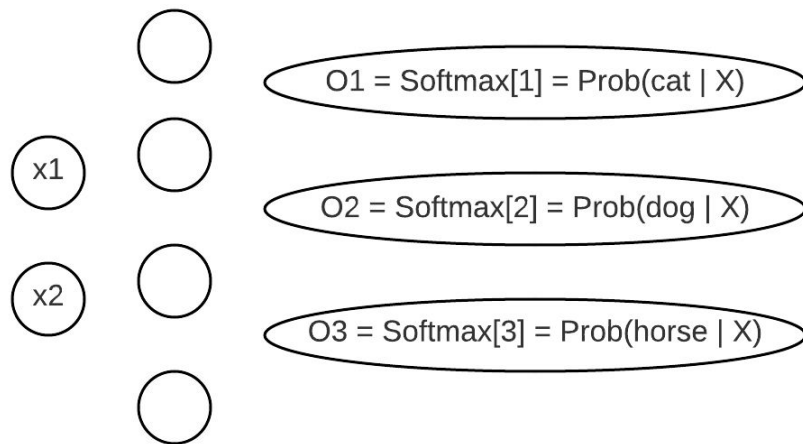
$$\mathcal{L} = \prod_{n=1}^N \prod_{i=1}^C \hat{y}_{ni}^{y_{ni}}$$

$$-\log(\mathcal{L}) = - \sum_{n=1}^N \sum_{i=1}^C y_{ni} \log(\hat{y}_{ni})$$

- This is the same as cross-entropy **but** over the whole  $N$  examples and move the negative to inside after the first sum!

# Probabilistic Formulation

- It is common to view the output of **NN for classification** as a conditional probability  $P(Y | X)$
- **Probability (class=cat | img)** can modeled as a **function** (e.g. NN, logistic reg) that transforms **input**  $X$  into the output **probability** of being e.g. a cat



# Discriminative Model

- A discriminative model **discriminate** (distinguish) between the **values** of the outcome
  - In classification (mainly) discriminate classes by finding their decision boundaries
- Examples
  - **Classifiers**: logistic regression, NN, SVM, KNN, etc
  - **Regressors**: Linear regression, NN, SVR, KNN, etc
- Modeled with  $P(Y|X)$ 
  - X is the input
    - E.g. image, house features, etc
  - Y is a class for classification and *continuous output variable* for regression
  - Learns input vs output relationship



# Generation

- What if I want the model to generate a new cat image not classify it?!
- What if a designer wants to describe a fancy scene and get an image?
- What if I have a text and wants to hear it with Trump voice?
- What if I want to chat with QA service like chatgpt or support services?
- What about code generation?
- This all requires the ability **to generate!**


Child joy person with blue eyes



# Generative Examples



## DALLE 2



## Midjourney



concerned teen with depression sitting on a bench inside a tunnel 8k ultra high quality, realistic, detailed, clear lines --ar 4:5 --v 5 --q 2 --s 750

# Generative Models

- Approaches that **explicitly or implicitly** model the distribution of (input, output)
  - Learn the **joint** probability distribution  $P(\mathbf{X}, \mathbf{Y})$  or the data distribution  $P(\mathbf{X})$  itself
- Then, we can use it to generate new **synthetic** samples
  - In chat, like Chatgpt, we can generate (predict) the next word in a sequence
    - LLM: Large Language models
      - LaMDA: Language Model for **Dialogue** Applications
    - Generative AI (GenAI): learns to generate a content (e.g. image, audio, language).
      - Pretrained (Foundation) models are fine-tuned for more apps
      - You may give a prompt / hallucination is a big challenge
- Compared to discriminative models: Requires more data, more computationally expensive but can be used for more than discrimination
  - many old models can be very hard to compute or get a weak model due to strong assumptions
  - Outliers negatively affects the resulting model



# Generative Models

- Assume we have input example  $X$  of  $n$  features
- And we have  $C$  labels
- We want to model  $P(C, X)$ , which can be expressed with the [chain rule](#) as
  - Chain rule:  $P(A, B) = P(B|A) P(A) = P(A|B) P(B)$

$$\begin{aligned} p(C_k, x_1, \dots, x_n) &= p(x_1, \dots, x_n, C_k) \\ &= p(x_1 \mid x_2, \dots, x_n, C_k) p(x_2, \dots, x_n, C_k) \\ &= p(x_1 \mid x_2, \dots, x_n, C_k) p(x_2 \mid x_3, \dots, x_n, C_k) p(x_3, \dots, x_n, C_k) \\ &= \dots \\ &= p(x_1 \mid x_2, \dots, x_n, C_k) p(x_2 \mid x_3, \dots, x_n, C_k) \cdots p(x_{n-1} \mid x_n, C_k) p(x_n \mid C_k) p(C_k) \end{aligned}$$

# Generative Models Approaches

- Gaussian Mixture Models
- Naive Bayes
- Hidden Markov Models
- Deep Learning
  - Generative Adversarial Networks (GANs)
  - Variational Autoencoders (VAEs)
  - [Diffusion](#) models (e.g. Stable Diffusion)
  - Autoregressive models (e.g. GPT: Generative Pre-trained Transformers)

# Question!

- You just enrolled in the CS department at Cairo University
- What might be the percentages of girls?
- In the morning in the first class, there were 39 man and 1 girl!
- One of the class students is coming from 200 meters far toward the class
- What is the probability of being a girl based on the data you have now?
- At 100 meters, seems the person is dressing a **skirt** ...maybe!
- Is the the probability of being a girl the same? Increasing? decreasing?
- At 50 meters, seems the person has a **very long hair**
- Is the the probability of being a girl the same? Increasing? decreasing?

# Prior, Likelihood and Posterior

- *Tip: ignore below beta calculations :)*
- Prior: Your **belief** before data.
  - So we assumed girls are very close to 50% as we see in many classes
  - It is our **prior** belief about the data without seeing the data e.g. Beta(50, 50)
- Likelihood: How the available **data** changes your **belief**.
  - Now, we realized current data is  $1/40 = 0.025$  of being a girl
  - But we feel skeptical being that far from what we think about the percentages
- Posterior: Your updated belief after considering the data.
  - To avoid data **extreme** conclusions, use your prior to update your belief
    - E.g. [Beta](#)(50+1,50+39)  $\Rightarrow E[\text{Beta}()] = (50 + 1) / (50 + 1 + 50 + 39) = 0.364$
  - Update the belief when you saw a skirt  $\Rightarrow$  **0.70**
  - Update the belief when you saw a long hair  $\Rightarrow$  **0.85**

# Bayes' Theorem

- **Bayes' Theorem:** [Update](#) the probability estimate with more data (evidence)
- If X is the input, Y is the output, K is one of the classes

$$\text{Posterior Probability} = \frac{\text{Likelihood} \times \text{Prior Probability}}{\text{Evidence}}$$

$$P(y = k | \mathbf{x}) = \frac{P(\mathbf{x} | y = k) P(y = k)}{P(\mathbf{x})}$$

$$P(\mathbf{x}) = \sum_k P(\mathbf{x} | y = k) P(y = k)$$

# Bayes' Theorem: Simple example

- We have an email that has the word “X = Free”. We want to know to what degree the email can be a “Y= Spam” because of this word
  - **Prior Probability  $P(Y)$** : Our **initial belief** that 20% of emails are spam
  - **Likelihood  $P(X|Y)$** : We got new information: email's text with **word free**  $\Rightarrow P(\text{free} | \text{spam})$ .
    - How probable the **observed data** is, **given** a set of **parameters** in the model:  **$P(X|\theta)$** 
      - If the email is really spam, what is the probability to see word free?
      - Observe It is  $P(\text{Input} | \text{Output})$  - **flipped**
  - **Evidence  $P(X)$** : The **overall probability** of observing word “free”, regardless spam or not
    - We loop on all possible classes to **marginalize**
  - **Posterior Probability  $P(Y|X)$** : The **updated** probability of being spam with word “free”
    - $P(\text{Output} | \text{Input})$ : is what we want for the conditional distribution
- Prior Probability (Spam) = **0.2** / Prior Probability (Not Spam) = 0.8
  - Sums to 1. For example our dataset has 200 spam emails and 800 non-spam emails
- Likelihood (“free” | Spam) = **0.7** / Likelihood (“free” | Not Spam) = 0.1

# Bayes' Theorem: Simple example

- Evidence (Marginal distribution / denominator)

- Iterate on possible classes: spam / !spam
- $P(X) = (0.2 \times 0.7) + (0.8 \times 0.1) = 0.14 + 0.08 = 0.22$

$$P(\mathbf{x}) = \sum_k P(\mathbf{x}|y = k)P(y = k)$$

- Posterior Probability (Spam | "free")

- $0.7 \times 0.2 / 0.22 = 0.64$

$$P(y = k|\mathbf{x}) = \frac{P(\mathbf{x}|y = k)P(y = k)}{P(\mathbf{x})}$$

- Overall


- Before we see a new email: priority probability of spam is 0.2
- When we saw a an email with word "free", the updated probability jumped to 0.64

- Generative Models

- $P(X, Y) = P(X | Y) P(Y) \Rightarrow$  Likelihood x Prior  $\Rightarrow$  computed in Bayes' theorem from data
- **Parameter Estimation** of the distribution may be done with [Maximum Likelihood Estimation](#)
- **Bayesian Inference**: parameters and predictions are treated as probability distributions

# From Discriminative to Generative

- Generative models model  $P(X,Y)$  or  $P(X|Y) + P(Y)$ .
  - In the Bayesian framework, these prior probabilities play a crucial role
- Discriminative models focus primarily on the likelihood, in the form  $P(Y|X)$ , while bayes rule enhance the likelihood with the prior





# Naive Bayes Classifier: A Generative Model

- Naive Bayes classifier is a **classification** algorithm: it has a **naive** assumption that all features are **mutually independent**, **conditional** on the **kth category**

$$p(x_i \mid x_{i+1}, \dots, x_n, C_k) = p(x_i \mid C_k).$$

Thus, the joint model can be expressed as

$$\begin{aligned} p(C_k \mid x_1, \dots, x_n) &= \propto p(C_k, x_1, \dots, x_n) \\ &\propto p(C_k) p(x_1 \mid C_k) p(x_2 \mid C_k) p(x_3 \mid C_k) \cdots \\ &\propto p(C_k) \prod_{i=1}^n p(x_i \mid C_k), \end{aligned}$$

$$p(C_k \mid x_1, \dots, x_n) = \frac{1}{Z} p(C_k) \prod_{i=1}^n p(x_i \mid C_k)$$

$$Z = p(\mathbf{x}) = \sum_k p(C_k) p(\mathbf{x} \mid C_k)$$

# Naive Bayes Classifier: Implementation

- Assume we have data

- X: (300, 5): 300 examples, each of 5 features
- Y: 3 classes 0, 1, 2

$$P(y = k|\mathbf{x}) = \frac{P(\mathbf{x}|y = k)P(y = k)}{P(\mathbf{x})}$$

- **Prior Probability**  $P(Y)$ : Let's use the frequency of each class
- Assume the underlying probability is gaussian
  - For each feature, compute its mean and std. This is its pdf (using MLE)
- **Likelihood**  $P(X[i]|\theta)$ : Given a sample  $\mathbf{x}$ , we can use gaussian pdf to evaluate
  - How probable this  $X[i]$  in this gaussian distribution
- **$P(\mathbf{X})$** : it is just a constant for a classifier
  - So, if we ignored, we have unscaled probabilities. We just needs the maximum
- Posterior Probability  $P(Y|\mathbf{X})$ : likelihood x prior

```
# Gaussian Probability Density Function
```

```
def gaussian_pdf(x, mean, std):  
    return (1 / (np.sqrt(2 * np.pi * std ** 2))) * \  
        np.exp(-((x - mean) ** 2) / (2 * std ** 2))
```

```
def get_data():  
    # Generate data that represents 3 classes  
    # for each class, generate 100 examples each of 5 features  
    # all data follows gaussian (mean, sigma) are provided  
    X0 = np.random.normal(2, 1, (100, 5))  
    X1 = np.random.normal(4, 1, (100, 5))  
    X2 = np.random.normal(6, 1, (100, 5))  
  
    # Combine into one dataset  
    X = np.vstack([X0, X1, X2]) # 300 x 5  
    y = np.array([0] * 100 + [1] * 100 + [2] * 100) # classes  
    # y: 100 0, then 100 1, then 100 2 for ground truth  
  
    # Create some test data  
    X_test = np.array([[1.5, 2, 2.2, 1.9, 2.1],  
                      [4.2, 3.8, 4.5, 3.9, 4.0],  
                      [6.1, 5.9, 6.0, 6.2, 6.1]])  
  
    return X, y, X_test
```

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

```
X, y, X_test = get_data()  
gnb = GNB()  
gnb.train(X, y)  
print(gnb.predict(X_test))
```

```
def train(self, X, y):
    self.classes = np.unique(y) # [0, 1, 2]
    self.means, self.stds, self.priors = {}, {}, {}

    for c in self.classes:
        X_c = X[y == c]
        self.means[c] = np.mean(X_c, axis=0)
        self.stds[c] = np.std(X_c, axis=0)
        self.priors[c] = len(X_c) / len(X) # frequency
```

- Notice, we collect all the data of cth class
- This represents the conditional dependency on the class  $P(x | c)$ 
  - Recall: conditional implies cut only those examples

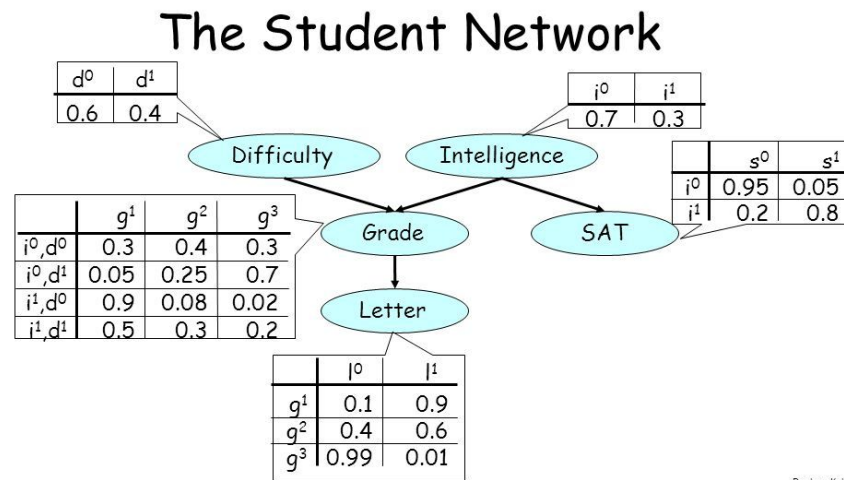
```
22 def predict(self, X):
23     num_samples = X.shape[0]
24     preds = np.zeros(num_samples)
25
26     for i in range(num_samples):
27         posteriors = {}
28
29         for c in self.classes:
30             # compute probability of each independent feature
31             props = gaussian_pdf(X[i, :], self.means[c], self.stds[c])
32             # multiply 5 props to get likelihood
33             likelihood = np.prod(props)
34             # compute posteriors, without constant Z
35             posteriors[c] = likelihood * self.priors[c]
36
37         # return the index of the max class value
38         preds[i] = max(posteriors, key=lambda k: posteriors[k])
39
40     return preds
```

# Frequentist vs. Bayesian inference

- Frequentist and Bayesian are two different approaches to **statistical inference**, and both have their merits and disadvantages.
  - used for making predictions, estimating parameters, and testing hypotheses
- Frequentist Inference
  - Views probability as the long-term frequency of events
  - Considers parameters to be fixed but unknown
    - MLE is common for estimating the parameters from the sampled data
    - MLE provides a point estimate without directly giving a measure of uncertainty
- Bayesian Inference
  - Views probability as a measure of belief.
  - Treats data as fixed once observed. Considers parameters to be random variables.
  - Allows for the inclusion of **prior** information.
- Good [article](#)

# Graphical Models

- Graphical models are a visual framework to represent complex systems and their probabilistic relationships using graphs
  - nodes represent random variables
  - edges represent probabilistic relationships between them
- Key professor Daphne Koller - course



# Probabilistic Modeling

- Why?
  - **Uncertainty** Quantification: Data (input/label), parameters  $P(\theta|D)$ , Model!
  - Flexibility ([un]supervised), Interpretability, Bayesian Updating
- Key Components: Random Variables, Parameters, Likelihood, Prior, Posterior Distribution
- Models
  - **Bayesian** Models: typically generative models, but can be discriminative  $P(\theta|D) = \frac{P(D|\theta) \times P(\theta)}{P(D)}$
  - **Graphical** Models: both generative and discriminative
    - Bayesian Networks, Markov Random Fields (MRFs), Factor Graphs, DBNs, Conditional Random Fields (CRFs)
  - **Time Series** Models: typically discriminative models (forecasting), but can be generative
    - Statistical (ARIMA, STL, ETS), Bayesian (BSTS, Gaussian Processes), Hidden Markov Models, State Space Models, Kalman Filters, Random Forest, Gradient Boosting, LSTM, Transformers (not all are probabilistic)



# Relevant Materials

- Likelihood
  - [Link](#) - [link](#) - [link](#) - [link](#) - [link](#) - [video](#)
- Probabilistic modeling
  - [Link](#)
- Bayesian or Frequentist
  - [Bayesian or Frequentist](#), Which Are You? By Michael I. Jordan (Part 1 of 2)
  - [Frequentism and Bayesianism](#): A Practical Introduction
  - Understanding [frequentist vs. Bayesian](#) inference

*“Acquire knowledge and impart it to the people.”*

*“Seek knowledge from the Cradle to the Grave.”*

