

Machine Learning

Padding and Stride

Mostafa S. Ibrahim

Teaching, Training and Coaching for more than a decade!

Artificial Intelligence & Computer Vision Researcher

PhD from Simon Fraser University - Canada

Bachelor / MSc from Cairo University - Egypt

Ex-(Software Engineer / ICPC World Finalist)



© 2023 All rights reserved.

Please do not reproduce or redistribute this work without permission from the author

Challenge at the Boundaries

- Assume we have a 3x3 input matrix and a kernel 3x3 for average ($c = 1/9$)
- What is the output matrix? Its size?
- What is the output for location `input[0][0]`?

1	2	3
4	5	6
7	8	9

c	c	c
c	c	c
c	c	c



Challenge at the Boundaries

- One solution is to only compute output for **valid** input locations!
- A valid pixel is one where the kernel completely inside the matrix
- There is only one valid pixel here at (1, 1) with input value 5
 - $(1+2+3+4+5+6+7+8+9) / 9 = 5$
- Given input DxD and kernel KxK, what is the **reduced** output matrix size?!

1	2	3
4	5	6
7	8	9

c	c	c
c	c	c
c	c	c



5

Challenge at the Boundaries

- Another approach is to pad the grid such that the output has the **same** dimensions as the input, then convolve on the valid pixels in padded grid

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0

c	c	c
c	c	c
c	c	c



1.3	2.3	1.8
3	5	3.7
2.7	4.3	3.1

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0

1st operation: $(0 + 0 + 0 + 0 + 1 + 2 + 0 + 4 + 5) / 9 = 1.3$

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0


6th operation: $(2+3+0+5+6+8+9+0) / 9 = 3.7$

Pytorch

- Pytorch provides us the flexibility of the 3 approaches for padding
 - Padding controls the amount of padding applied to the input.
 - String: **valid** \Rightarrow only valid pixels \Rightarrow grid size is reduced
 - String: **same** \Rightarrow pad with zeros such that output has the same input's resolution
 - Most common in practice
 - Specifically, pad with $(k-1)/2$ from each side
 - An int / a tuple of ints giving the amount of implicit padding applied on both sides.

PyTorch Kernel

- We can create a custom kernel. Here is a simple one for averaging
- It creates KxK matrix for averaging

```
class AverageConvolution(nn.Module):  
    def __init__(self, k = 3, padding = 'same'):  
        super(AverageConvolution, self).__init__()  
        self.padding = padding  
        self.filter = torch.ones(1, 1, k, k) / (k * k)  
  
    def forward(self, x):  
        return torch.conv2d(x, self.filter, padding=self.padding)
```

```
inp_d = 10
input = torch.arange(1, inp_d*inp_d+1).view(1, 1, inp_d, inp_d).float()

kernel = AverageConvolution(k=3, padding='same')
print(kernel(input).shape)      # torch.Size([1, 1, 10, 10])

# With valid, output is (D - K + 1) x (D - K + 1)
kernel = AverageConvolution(k=3, padding='valid')
print(kernel(input).shape)      # torch.Size([1, 1, 8, 8])

kernel = AverageConvolution(k=5, padding='valid')
print(kernel(input).shape)      # torch.Size([1, 1, 6, 6])

kernel = AverageConvolution(k=7, padding='valid')
print(kernel(input).shape)      # torch.Size([1, 1, 4, 4])

kernel = AverageConvolution(k=9, padding='valid')
print(kernel(input).shape)      # torch.Size([1, 1, 2, 2])

kernel = AverageConvolution(k=10, padding='valid')
print(kernel(input).shape)      # torch.Size([1, 1, 1, 1])
```


Stride

- What if we don't want a response for every pixel and would like to jump some in the middle?
 - For example, to reduce output size as in DNN
- The stride controls the **step size** at which the kernel is moved across the input data
 - A stride of 1 means the kernel moves on every pixel (as we did)
 - Most common stride

Stride

- Assume input is 6x6 and kernel = 3x3 with padding = valid (no padding)
- If the stride is S=3, this means jump 3 locations after each calculation
- Below 2 computations only per first 3 rows

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24
25	26	27	28	29	30
31	32	33	34	35	36

$$\begin{aligned} &(1+2+3+ \\ &7+8+9+ \\ &13+14+15) \\ &/ 9 = 8 \end{aligned}$$

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24
25	26	27	28	29	30
31	32	33	34	35	36

$$\begin{aligned} &(4+5+6+ \\ &10+11+12+ \\ &16+17+18) \\ &/ 9 = 11 \end{aligned}$$

Stride

- Similarly, we can evaluate the below 2 corners

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24
25	26	27	28	29	30
31	32	33	34	35	36

$$\begin{aligned} & (19+20+21+ \\ & 25+26+27+ \\ & 31+32+33) \\ & / 9 = 26 \end{aligned}$$

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24
25	26	27	28	29	30
31	32	33	34	35	36

$$\begin{aligned} & (22+23+24+ \\ & 28+29+30+ \\ & 34+35+36) \\ & / 9 = 29 \end{aligned}$$

Stride

- So overall, we get 2x2 matrix from stride=3
- Given input DxD, kernel KxK, No padding and stride S, the reduced matrix has a dimension $((D - K) / S) + 1$


1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24
25	26	27	28	29	30
31	32	33	34	35	36

c	c	c
c	c	c
c	c	c




8	11
26	29

Stride + Padding




no_padding_no_strides




no_padding_strides

Stride + Padding

Given input $D \times D$, kernel $K \times K$, padding P and stride S , the output matrix has a dimension of: $\lfloor (D - K + 2P) / S \rfloor + 1$



padding_no_strides




padding_strides

Dilated (Atrous) Convolution

- Each region where the kernel is extracting information is called a receptive field
- When we use a small 3x3 kernel, we see a very local region, so we lack more of the surrounding region
 - Later, we will have 3x3 weights for kernel. Bigger kernel are more computationally expensive.
- A meet in middle solution is to have a bigger view that works on a small subset of pixels by inserting gaps(dilations) between the kernel elements
- **Dilation Rate:** controls the spacing between the kernel points.
 - A dilation rate of 1 corresponds to the standard (non-dilated) convolution.
 - Larger dilation rates result in larger receptive fields.

Dilated (Atrous) Convolution

- This is a convolution operator with no stride or padding but dilation rate = 1
- It consists of 3x3 values computations, but extracting them from 5x5 region



torch.nn.Conv2d

- We now know many of pytorch [conv2d](#) parameters
- In future, explore **groups** parameter
 - controls the connections between inputs and outputs.
 - in_channels and out_channels must both be divisible by groups

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1,  
padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros',  
device=None, dtype=None) \[SOURCE\]
```

Applies a 2D convolution over an input signal composed of several input planes.

Relevant Materials

- A guide to [convolution arithmetic](#) for deep learning

“Acquire knowledge and impart it to the people.”

“Seek knowledge from the Cradle to the Grave.”

