

Machine Learning

Misc DNN Topics

Mostafa S. Ibrahim

Teaching, Training and Coaching for more than a decade!

Artificial Intelligence & Computer Vision Researcher

PhD from Simon Fraser University - Canada

Bachelor / MSc from Cairo University - Egypt

Ex-(Software Engineer / ICPC World Finalist)



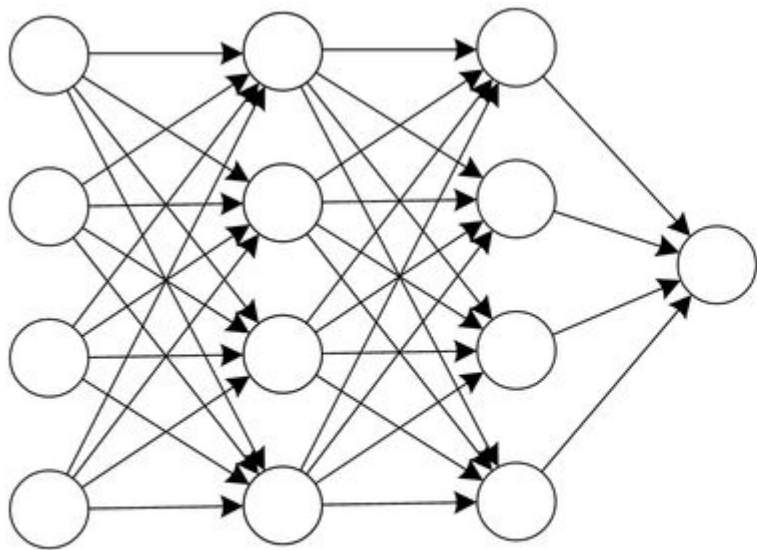
© 2023 All rights reserved.

Please do not reproduce or redistribute this work without permission from the author

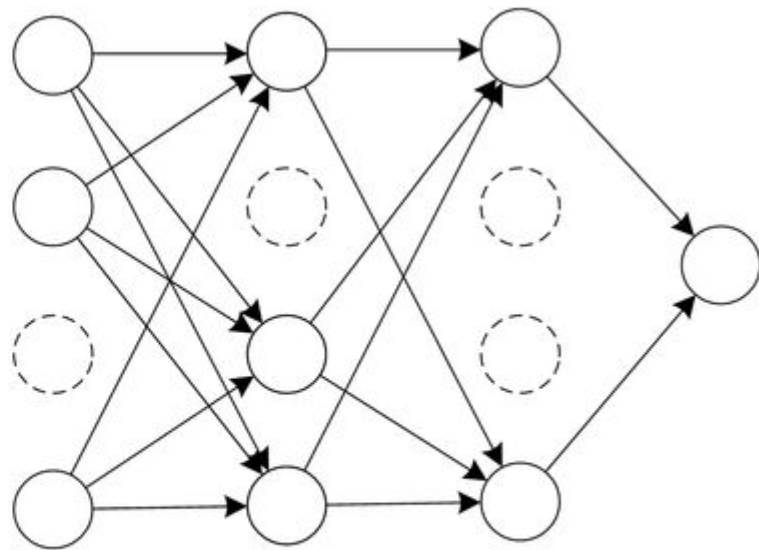
Dropout

- One of the great regularizers that fits overfitting in deep learning
- It works by randomly "dropping out" a number of neuron outputs in a layer during a training feedforward
 - This process prevents units from **co-adapting** too much to the data
 - The network learns the pattern rather than memorizing using its huge # of weights
 - **self.dropout = nn.Dropout(0.5)**
- It is applied on training only. On testing, all nodes are applied
 - `model.eval()` is critical
- Where to apply?
 - Mainly In front of linear layers with maybe dropout = 0.5
 - You may explore after some pooling layers with dropout = 0.1 / 0.2
- [Video](#)

Dropout



(a) Standard Neural Network



(b) Network after Dropout

Batchnorm (2015)

- Observation: distribution of each layer's inputs changes during training
 - Reason: as the parameters of the previous layers change
 - Consequence: slow down the training process and make it harder for the network to converge
- Solution: **Normalize** layer's input + learn scale/shift **parameters**
 - In **train**: Use batch mean and batch std
 - In **inference**: uses the **entire** training set's **moving average** of the mean and variance
 - Again needs model.eval()
 - Addition side effects: Regularization Effect / Faster Convergence
 - Cons: Dependency on Batch Size / not straightforward for dynamic RNNs
 - [Alternatives](#): Layer Normalization, Instance Normalization, and Group Normalization
- Where to apply?
 - After convolutional layers / fully connected layers but before activation function (debate)
 - Tip: if a batchnorm after conv layer, don't use bias term in conv (redundant)
 - Resnet blocks

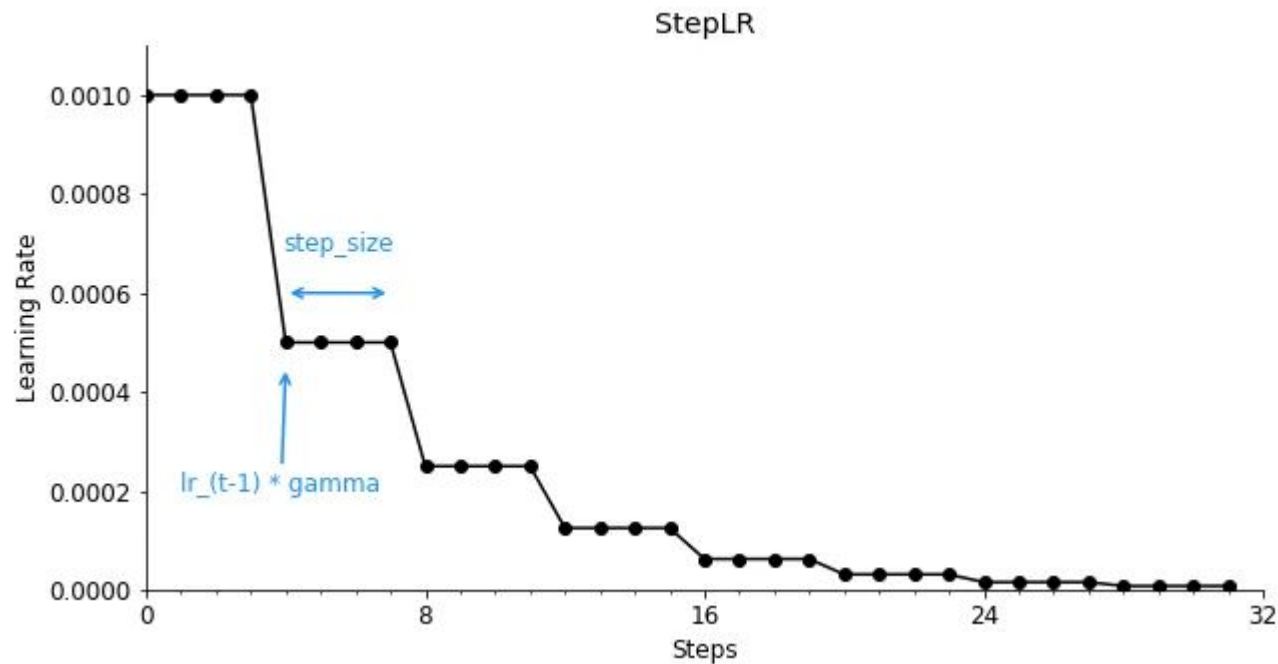
AdamW Optimizer [high level]

- AdamW is an improvement from Adam Optimizer
 - AdamW fixes an implementation mistake about the weight decay
 - It works on really many datasets. Consider as your starting point
- You can think of it as combination of **SGD with momentum** and **RMSProp** optimizers
- Adam maintains **two moving averages** for each parameter
 - At each step, Adam computes **adaptive learning rates for each parameter**
- Videos: Exponentially Weighted [Averages](#) / Gradient Descent With [Momentum](#) / [Adam](#) / [Overall](#)

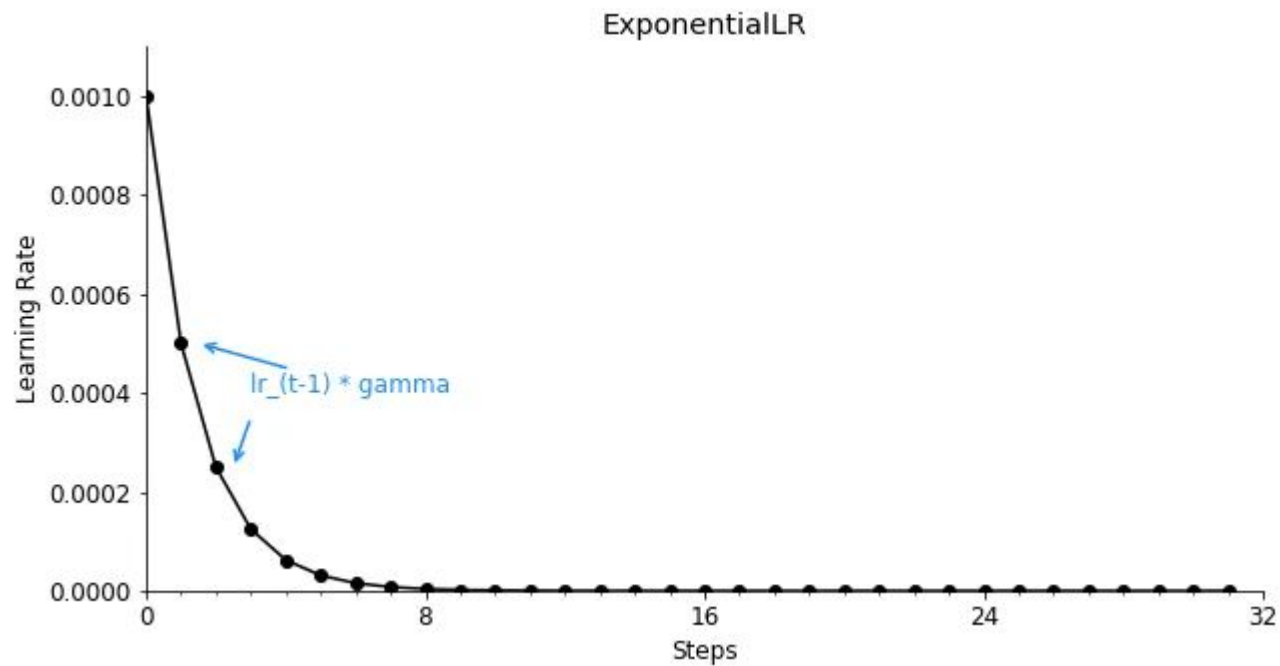
Learning Rate Schedulers [[Article](#)]

- We don't just fix one learning rate. We need it to vary during the training
- There are many strategies for that
 - **StepLR**: One of the easiest/oldest ways
 - Decays the learning rate of each parameter group by gamma every step_size epochs.
 - `scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=30, gamma=0.1)`
 - Then, the gamma and step_size are hyperparameters!
 - Gamma like 1/10 or 1/2 are common choices
 - Check the loss curve to decide a good place for the step_size
 - MultiStepLR gives more variable milestones [30,80]
 - ExponentialLR: Decays by gamma **every epoch** [common]
 - ReduceLROnPlateau: when a metric has stopped improving
 - CosineAnnealingLR
 - OneCycleLR: [read](#)

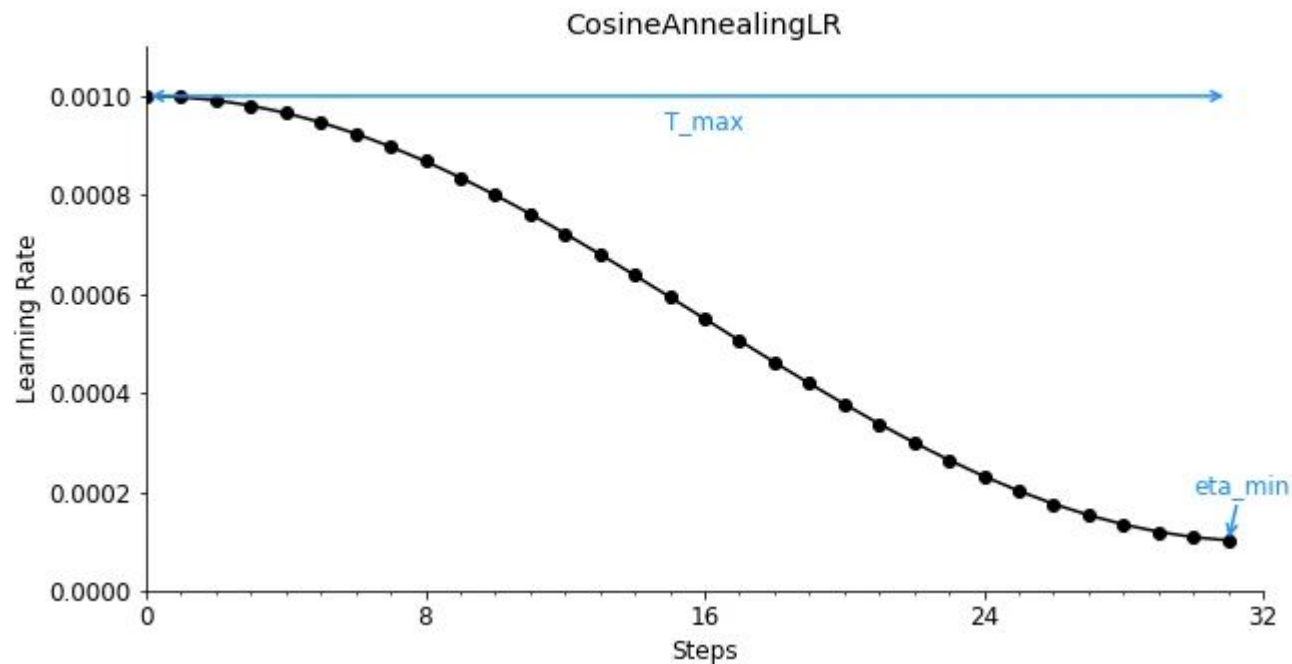
Learning Rate Schedulers



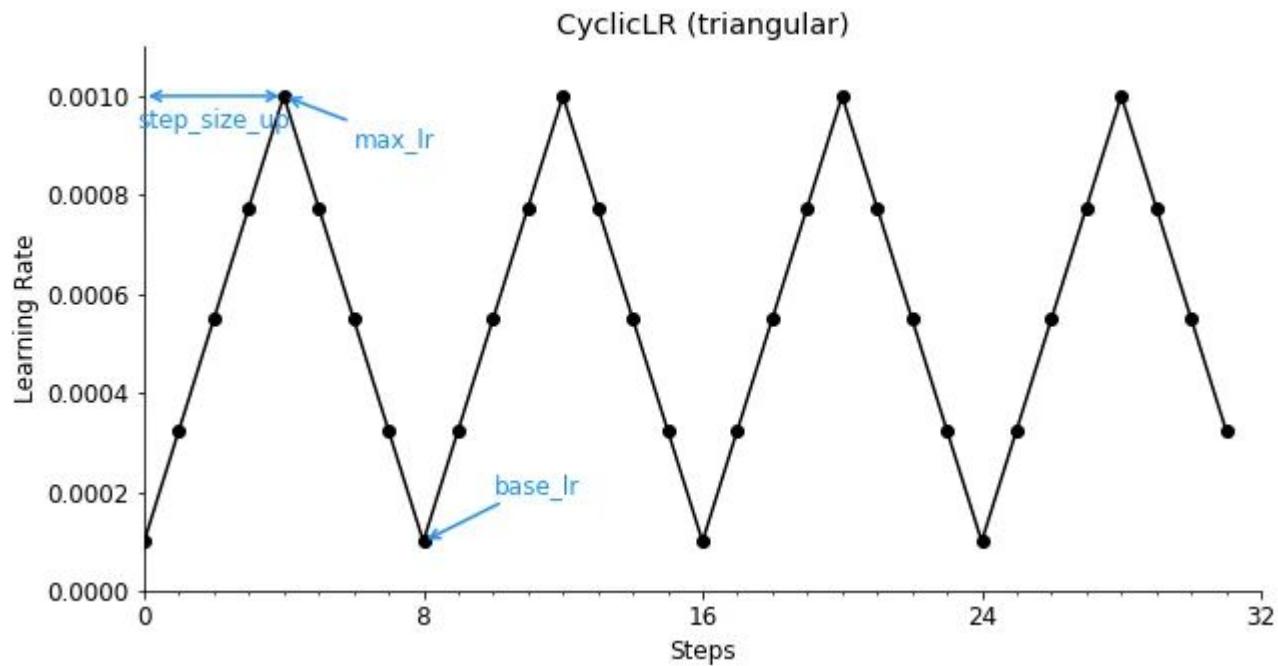
Learning Rate Schedulers



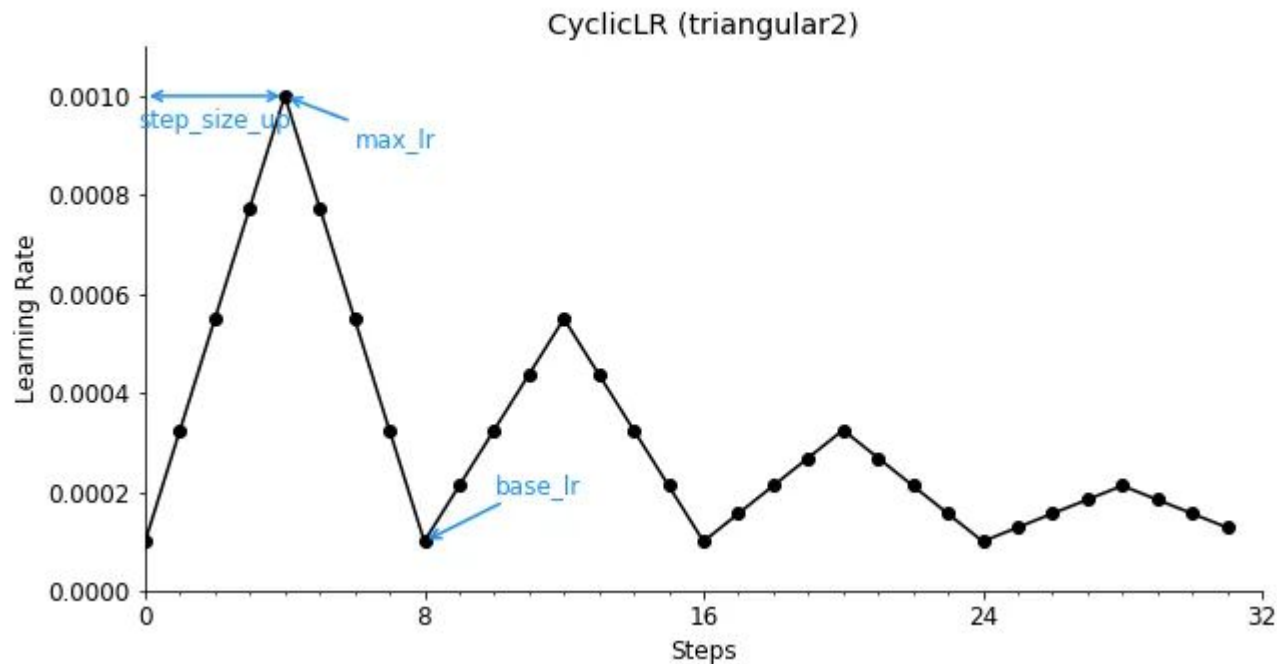
Learning Rate Schedulers



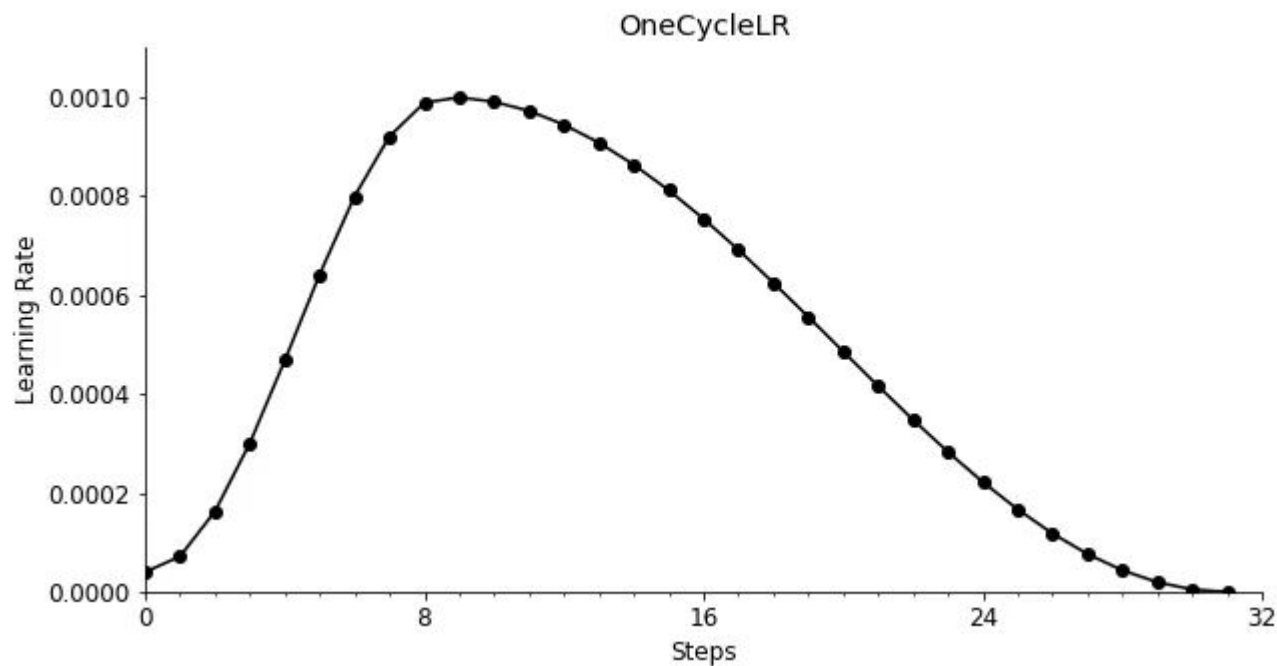
Learning Rate Schedulers



Learning Rate Schedulers



Learning Rate Schedulers



Vanishing Gradients

- During backpropagation in deep networks the gradients become very small the more we move backward in the hidden layers
 - As a result no or very slow learning as the gradient signal is too weak
 - The opposite is called exploding gradients
- Causes
 - **Improper initialization:** In feedforward, we are **multiplying weights** from the matrices
 - **Activation Functions:** In backpropagation, with chain rule, we **multiply gradients**
 - Sigmoid and tanh activations become extremely small for small/large inputs
 - The deeper the network, the harder to train due to these 2 reasons
 - Similarly long sequences in RNNs

Vanishing Gradients

- A good solution must tackle all these issues
 - **Weight Initialization:** Using He initialization or Xavier initialization
 - **Activation Functions:** Relu (+ve range derivative = 1) / non-saturating function
 - A saturating function: approach a plateau as the input grows large (sigmoid/tanh)
 - **Skip connections** (like in ResNets) allow gradients to bypass layers
 - **Batch Normalization:** Normalizing the input of each layer to have mean=0, variance=1
 - **Gradient Clipping:** Put a range to clip
- For exploding gradients, in addition to above solutions, **weight regularization** and **lower learning rate** are used

Network Initialization

- Initializing the network is a very critical component in training DNNs
- Early days, people realized some lessons
 - Very small weights leads to vanishing gradients and large ones to exploding gradients
 - Weights can't be the same constant. They should be different and have variance
 - Normal and Uniform distribution are good ways
- However, this was not that enough with deep learning
 - Weight initialization methods need to be compatible with the choice of an **activation function**, mismatch can potentially affect training negatively.
 - Just random initialize the weight may cause vanishing/exploding gradient

Network Initialization

- There are 2 major techniques nowadays to tackle deep learning
- **Xavier**/Glorot initialization [[paper](#)]
 - “Adjust the scale of the initial weights based on the number of **input and output** neurons in a way that aims to keep the **variance** of the activations **constant** across layers.”
 - Some say it is good for tanh/sigmoid activations, but TF uses as default / good for DNN
- **Kaiming**/He initialization
 - Designed specifically for **Deep Networks with RELU**: paper [Delving Deep into Rectifiers](#)
 - Most PyTorch layers use **Kaiming with Uniform distribution** for most of the layers
 - Such as Conv, Linear and RNN
 - However, the bias is sampled from the uniform distribution
 - These distributions are sampled from a **specific range** based on tensor input/output

Network Initialization: Custom

```
def init_weights(m):  
    if isinstance(m, nn.Conv2d):  
        init.xavier_uniform_(m.weight)  
        if m.bias is not None:  
            init.zeros_(m.bias)  
    elif isinstance(m, nn.BatchNorm2d):  
        init.constant_(m.weight, 1)  
        init.constant_(m.bias, 0)  
    elif isinstance(m, nn.Linear):  
        init.kaiming_normal_(m.weight)  
        if m.bias is not None:  
            init.zeros_(m.bias)  
  
model = nn.Sequential(  
    nn.Conv2d(1, 20, 5),  
    nn.ReLU(),  
    nn.Conv2d(20, 64, 5),  
    nn.ReLU(),  
    nn.Linear(64, 10)  
)  
  
model.apply(init_weights)
```

Tensorboard Visualization

- TensorBoard is a **visualization toolkit** (web app) for machine learning experimentation inspection developed by the TensorFlow team.
 - Adopted in PyTorch and other machine learning frameworks
- You can
 - Track and visualize **metrics** such as loss and accuracy
 - Visualize the model graph
 - Display **images**, text, and audio data samples
 - View the **distribution** of weights, biases, or other tensors as they change over time.
 - Track hyperparameter tuning sessions using the HPparams dashboard.
 - `from tensorboard.plugins.hparams import api as hp / writer.add_hparams()`
- Install from: **pip install tensorboard**
- Then build and visualize web from: **tensorboard --logdir=<Path>**

TensorBoard

TIME SERIES

SCALARS

IMAGES

🔍 Filter runs (regex)



Run ↑

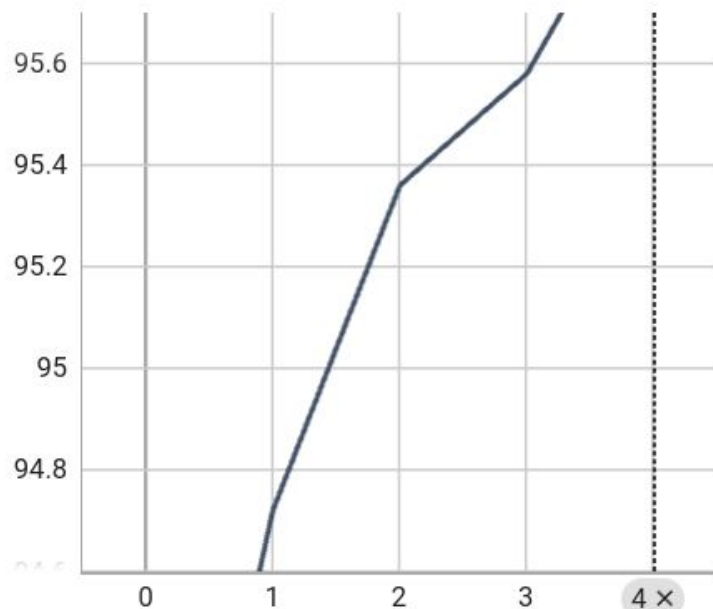


🔍 Filter tags (regex)

four_mnist_images



validation_accuracy



Run ↑

Value

Step

Relative



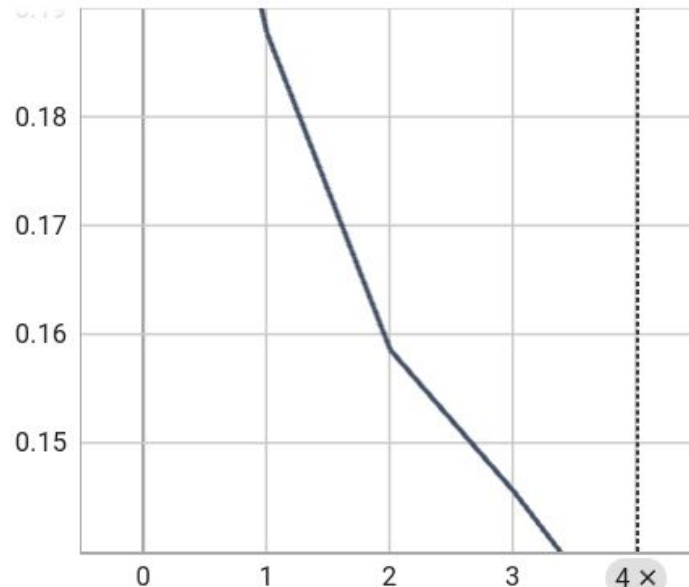
.

96.02

4

24.92 sec

validation_loss



Run ↑

Value

Step

Relative



.

0.1308

4

24.92 sec

```
writer = SummaryWriter('runs/mnist_experiment_1')
```

```
# Log the running loss averaged per mini-batch
```

```
writer.add_scalar('training_loss', running_loss / 100, epoch * len(train_loader) + i)
```

```
# Log a random batch of images
```

```
img_grid = torchvision.utils.make_grid(inputs[:4].cpu().data)
```

```
writer.add_image('four_mnist_images', img_grid, epoch * len(train_loader) + i)
```

```
# Log validation loss and accuracy
```

```
writer.add_scalar('validation_loss', avg_val_loss, epoch)
```

```
writer.add_scalar('validation_accuracy', val_accuracy, epoch)
```

```
print('Finished Training')
```

```
# Call the training loop
```

```
train(num_epochs=5)
```

```
# Close the TensorBoard writer
```

```
writer.close()
```

- See, run and visualize the full code

What you need in a good project?

- Data Loader
 - Always start with data loader and verify the data properly
 - Prepare sample mode (for fast check / debugging) and full mode loadings
- Configuration
 - Any hyperparameter or model choices must be from a configuration file
- Versioning
 - You must be able to know which code/version/data generated these results
 - For example, git code version + copy config files at minimum
- Logger
 - You must log every important thing such as model layers, data path, configuration content
- Events
 - Save events for later tensorboard visualization
- Multi-GPU support / train/eval round strategy

Google Colab

- Colab is a free, cloud-based service that supports ML
 - Write and execute Python in your browser
 - GPUs and TPUs for intensive computational tasks.
 - Integration with Google Drive: save and share your work
 - Supports TensorFlow, PyTorch, Keras, and OpenCV
 - Easy to use for collaboration
 - No setup

Relevant Materials

- [Weight Initialization](#) in Neural Networks
- [Weight Initialization](#) Techniques-What best works for you
 - Focus on the text not the equations
- What is [default](#) weight and bias initialization in PyTorch?

“Acquire knowledge and impart it to the people.”

“Seek knowledge from the Cradle to the Grave.”

